

# Lecture 3

## Introduction to optimization: Gradient descent

GEOL 4397: Data analytics and machine learning for geoscientists

Jiajia Sun, Ph.D.

Jan. 24<sup>th</sup>, 2019

UNIVERSITY of  
**HOUSTON**

YOU ARE THE PRIDE


EARTH AND ATMOSPHERIC SCIENCES



# Today's agenda

- Motivation
- Concept: gradient
- Gradient descent
  - Stochastic gradient descent
  - Mini-batch gradient descent
- Learning rate

# Motivation



Each point represents one training example/instance.

Figure from Aurelien Geron's ML book, page 19

# General approach to learning/training

defining a cost function  
&  
minimizing it

# Minimization

- Cost function measures how bad a candidate model is

$$\min J(\theta_0, \theta_1) = \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$

- Learning/training = Minimization = Optimization
- **Optimization**: finding optimal parameter values that minimize a cost function

# Matrix-vector form

$$J(\theta) = \sum_{i=1}^M (\boldsymbol{\theta}^T \mathbf{x}^{(i)} - y^{(i)})^2$$



$$J(\theta) = \|\mathbf{X}\boldsymbol{\theta} - \mathbf{y}\|^2$$

$$\mathbf{X} = \begin{bmatrix} (\mathbf{x}^{(1)})^T \\ \vdots \\ (\mathbf{x}^{(M)})^T \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(M)} \end{bmatrix}$$

$M \times (N+1)$                        $M \times 1$

# Analytical solution

- Minimize:

$$J(\theta) = \|X\theta - y\|^2$$

$$\tilde{\theta} = (X^T X)^{-1} (X^T y)$$

# Analytical solution

- Minimize:

$$J(\theta) = \|X\theta - y\|^2$$

- Normal equation method

$$\tilde{\theta} = (X^T X)^{-1} (X^T y)$$



# Analytical solution

- Minimize:

$$J(\theta) = \|X\theta - y\|^2$$

- Normal equation method

$$\tilde{\theta} = (X^T X)^{-1} (X^T y)$$

```
theta = np.matmul(np.linalg.inv(np.matmul(X.T,X)), np.matmul(X.T,y))
```

# To derive normal equation (optional)

- <http://www.programming-techniques.com/2013/12/gradient-descent-versus-normal-equation.html>
- <http://cs229.stanford.edu/notes/cs229-notes1.pdf>  
page 8-11

# Problem with normal equation method

$$\tilde{\theta} = (X^T X)^{-1} (X^T y)$$

# Problem with normal equation method

$$\tilde{\theta} = \underbrace{(X^T X)^{-1}}_{(N+1) \times (N+1)} (X^T y)$$

## Computational cost

- increases **linearly** with **M** (# of instances)

# Problem with normal equation method

$$\tilde{\theta} = \underbrace{(X^T X)^{-1}}_{(N+1) \times (N+1)} (X^T \mathbf{y})$$

## Computational cost

- increases **linearly** with **M** (# of instances)
- Increases **cubically\*** with **N** (# of features)

\*Strictly speaking, computational complexity is  $O(N^{2.4})$  to  $O(N^3)$ . If we double the number of features, the computation time increase by  $2^{2.4} = 5.3$  to  $2^3 = 8$  times

# Problem with normal equation method

$$\tilde{\theta} = \underbrace{(X^T X)^{-1}}_{(N+1) \times (N+1)} (X^T y)$$

## Computational cost

- increases **linearly** with **M** (# of instances)
- Increases **cubically\*** with **N** (# of features)
- For **big data, computationally very expensive**

\*Strictly speaking, computational complexity is  $O(N^{2.4})$  to  $O(N^3)$ . If we double the number of features, the computation time increase by  $2^{2.4} = 5.3$  to  $2^3 = 8$  times

# Problem with normal equation method

$$\tilde{\theta} = \underbrace{(X^T X)^{-1}}_{(N+1) \times (N+1)} (X^T y)$$

## Computational cost

- increases **linearly** with **M** (# of instances)
- Increases **cubically\*** with **N** (# of features)
- For **big data, computationally very expensive**

Think about the **life satisfaction** problem, we only used one feature, i.e., **GDP per capita**. **What other features could we use?**

\*Strictly speaking, computational complexity is  $O(N^{2.4})$  to  $O(N^3)$ . If we double the number of features, the computation time increase by  $2^{2.4} = 5.3$  to  $2^3 = 8$  times

# Problem with normal equation method

$$\tilde{\theta} = \underbrace{(X^T X)^{-1}}_{(N+1) \times (N+1)} (X^T y)$$

## Computational cost

- increases **linearly** with **M** (# of instances)
- Increases **cubically\*** with **N** (# of features)
- For **big data, computationally very expensive**

For nonlinear optimization, normal equation does not even exist.

\*Strictly speaking, computational complexity is  $O(N^{2.4})$  to  $O(N^3)$ . If we double the number of features, the computation time increase by  $2^{2.4} = 5.3$  to  $2^3 = 8$  times




# Gradient descent

- Computationally less demanding
- Generally applicable to both linear and nonlinear optimization\*

\*so long as gradient can be calculated.

# What is gradient?

- Let us first recall what is derivative.



Picture taken from <https://en.wikipedia.org/wiki/Derivative>

# From derivative to gradient

- Let us consider a function  $f(x, y)$
- Two partial derivatives

$$\frac{\partial f}{\partial x}$$

$$\frac{\partial f}{\partial y}$$

# Gradient

- Gradient of a function  $f(x, y)$  is defined as

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$$

# More on gradient

- It is a **vector**
- Therefore, it has **direction** and **magnitude**
- Its **direction** points in the direction of the greatest rate of increase (i.e., direction of maximum increase) of the function
- Its **magnitude** is the slope of the graph of the function (i.e., **the rate of increase**) in that direction

# Put gradient in context

- Imagine you are standing on a hillside. Look all around you, and find the **direction of steepest ascent**.

# Put gradient in context

- Imagine you are standing on a hillside. Look all around you, and find the **direction of steepest ascent**.
- That is **the direction of the gradient**.

# Put gradient in context

- Imagine you are standing on a hillside. Look all around you, and find the direction of steepest ascent.
- That is the direction of the gradient.
- Now measure **the slope in that direction** (rise over run)

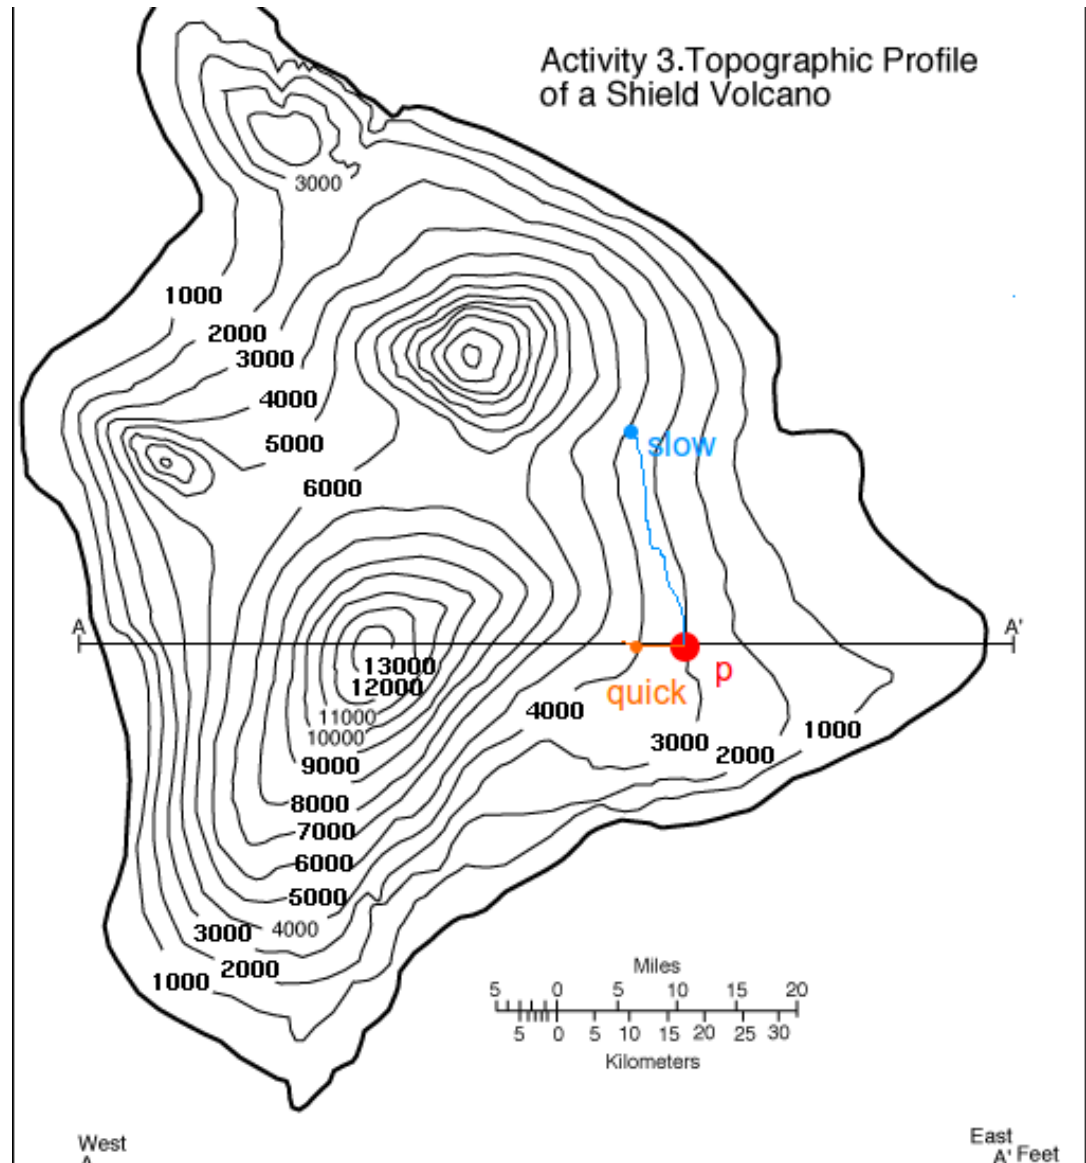


# Put gradient in context

- Imagine you are standing on a hillside. Look all around you, and find the direction of steepest ascent.
- That is the direction of the gradient.
- Now measure **the slope in that direction** (rise over run)
- That is the **magnitude** of the gradient.
- Here, the function is the height of hill (as a function of positions).

# Understanding gradient


- Consider the topography as a 2D function  $f(x, y)$
- The gradient direction tells you the fastest way up



Picture taken from <https://mathoverflow.net/questions/1977/why-is-the-gradient-normal>

# Gradient in the context of optimization

- Optimization problem is often posed as a minimization problem



- We want to find where the minimum of a cost function is.

Picture taken from Andrew Ng's Machine Learning class on Coursera.org

# Gradient descent algorithm

- Given initial values  $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$
- While (not convergence):

$$\boldsymbol{\theta}^{(j)} = \boldsymbol{\theta}^{(j-1)} - \alpha \nabla J(\boldsymbol{\theta}^{(j-1)})$$

# Gradient descent algorithm

- Given initial values  $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$
- While (not convergence):

$$\theta_0^{(j)} = \theta_0^{(j-1)} - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0^{(j-1)}, \theta_1^{(j-1)})$$
$$\theta_1^{(j)} = \theta_1^{(j-1)} - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0^{(j-1)}, \theta_1^{(j-1)})$$

# Gradient descent algorithm for linear regression

- Given initial values  $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$
- While (not convergence):

$$\theta_0^{(j)} = \theta_0^{(j-1)} - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0^{(j-1)}, \theta_1^{(j-1)})$$
$$\theta_1^{(j)} = \theta_1^{(j-1)} - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0^{(j-1)}, \theta_1^{(j-1)})$$

$$\min J(\theta_0, \theta_1) = \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$

# Gradient descent algorithm for linear regression

- Given initial values  $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$
- While (not convergence):

$$\theta_0^{(j)} = \theta_0^{(j-1)} - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0^{(j-1)}, \theta_1^{(j-1)})$$
$$\theta_1^{(j)} = \theta_1^{(j-1)} - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0^{(j-1)}, \theta_1^{(j-1)})$$

$$\min J(\theta_0, \theta_1) = \frac{1}{2M} \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$

# Gradient descent algorithm for linear regression

- Given initial values  $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$
- While (not convergence):

$$\theta_0^{(j)} = \theta_0^{(j-1)} - \alpha \frac{1}{M} \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})$$


$$\theta_1^{(j)} = \theta_1^{(j-1)} - \alpha \frac{1}{M} \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) x^{(i)}$$



Slides 36-44 are from Andrew Ng's machine learning course on [coursera.org](https://www.coursera.org)


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$


(function of the parameters  $\theta_0, \theta_1$ )






$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )



# Gradient descent algorithm

- Given initial values  $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$
- While (not convergence):

$$\theta_0^{(j)} = \theta_0^{(j-1)} - \alpha \frac{1}{M} \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})$$

$$\theta_1^{(j)} = \theta_1^{(j-1)} - \alpha \frac{1}{M} \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) x^{(i)}$$

# Observation

- Each step of gradient descent uses ALL the training examples.

# Batch gradient descent

- Each step of gradient descent uses ALL the training examples.

# Problem with batch gradient descent

- When the number of training data is huge, say,  $M = 300,000,000$ , batch gradient descent becomes very slow.

# Problem with batch gradient descent

- When the number of training data is huge, say,  $M = 300,000,000$ , batch gradient descent becomes very slow.

In 2017,

67 million Instagram posts uploaded each day!

656 million Tweets were generated each day!

4.3 billion Facebook messages posted daily!

<https://blog.microfocus.com/how-much-data-is-created-on-the-internet-each-day/>

# Gradient descent algorithm

- Given initial values  $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$
- While (not convergence):

$$\theta_0^{(j)} = \theta_0^{(j-1)} - \alpha \frac{1}{M} \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})$$

$$\theta_1^{(j)} = \theta_1^{(j-1)} - \alpha \frac{1}{M} \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) x^{(i)}$$



- Given initial values  $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$

$$\theta_0 = \theta_0 - \alpha (\theta_0 + \theta_1 x^{(i)} - y^{(i)})$$

$$\theta_1 = \theta_1 - \alpha (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) x^{(i)}$$

- Given initial values  $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$

randomly shuffle indices  $[0, 1, 2, \dots, M-1]$

For  $i$  in shuffled indices {

$$\theta_0 = \theta_0 - \alpha (\theta_0 + \theta_1 x^{(i)} - y^{(i)})$$

$$\theta_1 = \theta_1 - \alpha (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) x^{(i)}$$

}

- Given initial values  $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$

randomly shuffle indices  $[0, 1, 2, \dots, M-1]$

For  $i$  in shuffled indices {

$$\theta_0 = \theta_0 - \alpha (\theta_0 + \theta_1 x^{(i)} - y^{(i)})$$

$$\theta_1 = \theta_1 - \alpha (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) x^{(i)}$$

}

one pass  
one epoch

# Stochastic gradient descent

- Given initial values  $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$

Repeat {

randomly shuffle indices [0,1,2, ..., M-1]

For  $i$  in shuffled indices {

$$\theta_0 = \theta_0 - \alpha (\theta_0 + \theta_1 x^{(i)} - y^{(i)})$$

$$\theta_1 = \theta_1 - \alpha (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) x^{(i)}$$

}



}

one pass  
one epoch

# SGD vs BGD


- SGD much **faster**
- Search path very **irregular**
- Cost function bounces up and down, decreasing only on average
- Over time, it ends up close to minimum, but never settles down.

# SGD vs BGD



Picture taken from  
<https://www.cs.cmu.edu/~yuxiangw/docs/SSGD.pdf>

# SGD cost function



<https://upload.wikimedia.org/wikipedia/commons/f/f3/Stogra.png>

# Mini-batch gradient descent

- Mini-batch uses a small number ( $1 < \# < M$ ) of training examples to update model parameters



# Batch vs stochastic vs Mini-batch





Figure from Aurelien Geron's ML book, page 120

# Learning rate

$$\theta_1^{(j)} = \theta_1^{(j-1)} - \alpha \frac{\partial}{\partial \theta_1} J(\theta_1^{(j-1)})$$

If  $\alpha$  is too small, gradient descent becomes very slow.



# Learning rate

$$\theta_1^{(j)} = \theta_1^{(j-1)} - \alpha \frac{\partial}{\partial \theta_1} J(\theta_1^{(j-1)})$$

If  $\alpha$  is too large, you might overshoot the minimum. It may fail to converge, sometimes even diverge.

