

Lecture 10

Review

GEOL 4397: Data analytics and machine learning for geoscientists

Jiajia Sun, Ph.D.
March 19th, 2019

UNIVERSITY of
HOUSTON

YOU ARE THE PRIDE

EARTH AND ATMOSPHERIC SCIENCES



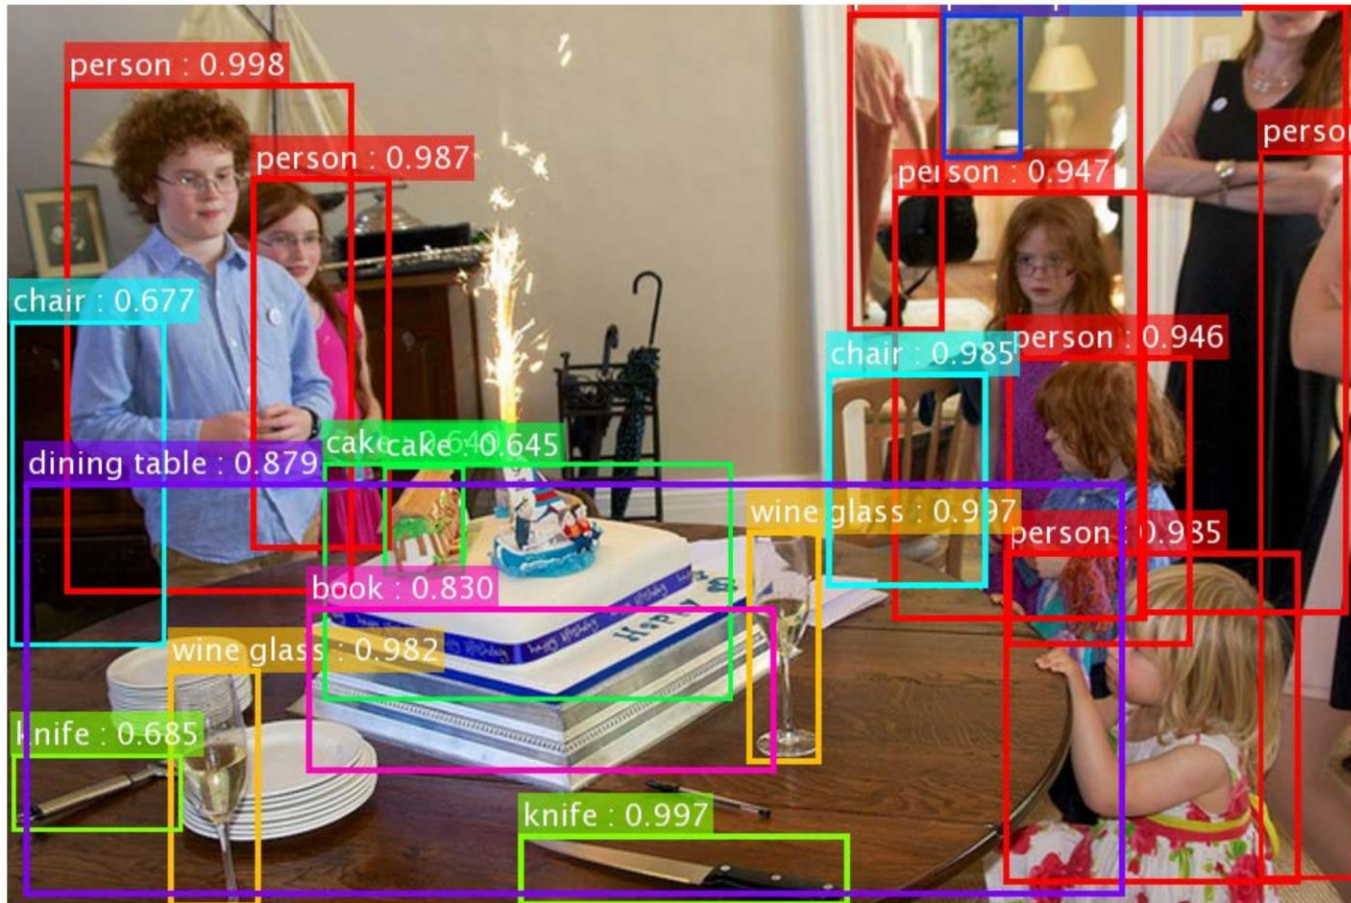
Announcement

- The lab report on ensemble learning due at 5:30 pm, March 21st.
- Exam on March 21st from 5:30 to 6:50 pm.

Outline

- Machine learning basics
 - Example applications
 - Definitions
- Training
 - Cost function
 - Optimization
- Optimization algorithms
 - Batch gradient descent
 - Stochastic gradient descent
 - Mini-batch gradient descent
- Types of machine learning
 - Supervised vs. unsupervised
 - Regression vs. Classification
- Overfit vs Underfit
 - Diagnose
 - Remedy


Object detection



ResNet applied to COCO dataset.


Source: He et al., Deep residual learning for image recognition, CVPR, 2016

Real time object detection



Video online: <https://www.youtube.com/watch?v=zZe27JYi8Y>

Self-driving car



Source: Google

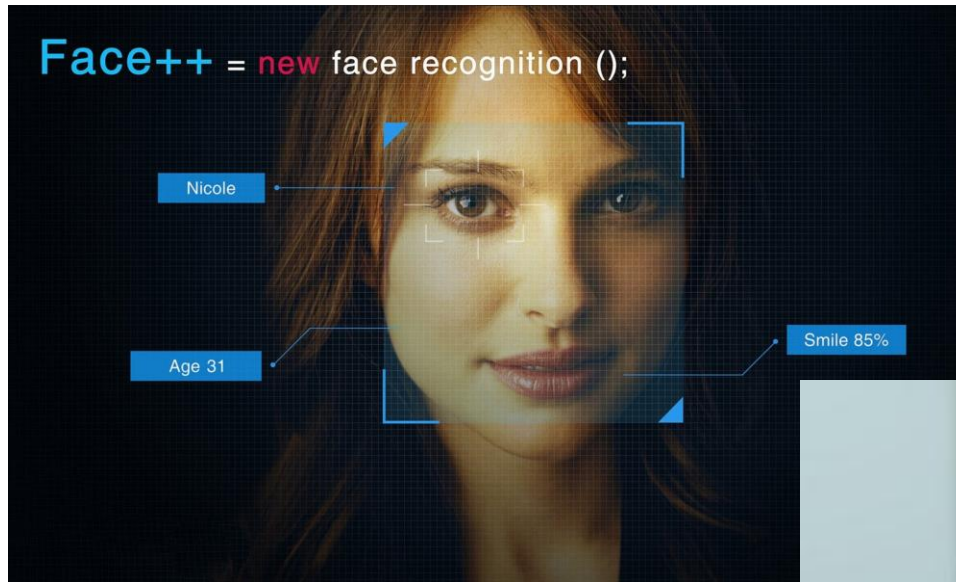
Raoul Rañoa / @latimesgraphics

Voice recognition




Source: <https://biostore.co.uk/company/news-articles/voice-recognition-biometrics-making-noise/>

Face recognition



Source: <https://www.pinterest.com/pin/135600638759424941/?lp=true>

Hand written digit recognition



MNIST data set

Spam filter



Fraud detection



Source: <http://tsigroup.com>

Recommender system

Hands-On Machine Learning with Scikit-Learn & TensorFlow

by Aurélien Géron (Author)
★★★★★ 134 reviews
#1 Best Seller in Artificial Intelligence

Frequently bought together



Total price: \$122.37


Add all three to Cart

Add all three to List


These items are shipped from and sold by different sellers. Show details

- ☒ This item: Python Machine Learning: Machine Learning and Deep Learning with Python, scikit-learn, and... by Sebastian F
- ☒ Hands-On Machine Learning with Scikit-Learn and TensorFlow: Concepts, Tools, and Techniques to Build... by Aurélien G
- ☒ Deep Learning (Adaptive Computation and Machine Learning series) by Ian Goodfellow Hardcover \$55.38


Customers who bought this item also bought




ISBN-13: 978-149196229
ISBN-10: 1491962291
Why is ISBN important?




Hands-On Machine Learning with Scikit-Learn and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems
by Aurélien Géron
★★★★★ 115 reviews
#1 Best Seller in Artificial Intelligence
Paperback
\$31.00 ✓prime



Artificial Intelligence with Python: A Comprehensive Guide to Building...
by Prateek Joshi
★★★★★ 6 reviews
Paperback
\$44.99 ✓prime



Deep Learning (Adaptive Computation and Machine Learning series)
by Ian Goodfellow
★★★★★ 116 reviews
Hardcover
\$55.38 ✓prime



Python for Data Analysis: Data Wrangling with Pandas, NumPy, and...
by Wes McKinney
★★★★★ 4 reviews
Paperback
\$34.82 ✓prime

s, and Techniques to Build

✓prime \$25.95

✓prime \$27.99

Price: \$49.99 Save: \$22.00 (44%)
35 New from \$27.89

1

Add to Cart

Turn on 1-Click ordering

Ship to:

Yu Zhang- Norman - 73072

68 used & new from \$19.01

Go game



Image credit: Nature



Image credit: theverge.com

What is machine learning?

- the field of study that gives computers the ability to **learn from data** (e.g., discovering **patterns** and **relations** among input data), and **make predictions**.

Outline

- Machine learning basics
 - Example applications
 - Definitions
- Training
 - Cost function
 - Optimization
- Optimization algorithms
 - Batch gradient descent
 - Stochastic gradient descent
 - Mini-batch gradient descent
- Types of machine learning
 - Supervised vs. unsupervised
 - Regression vs. Classification
- Overfit vs Underfit
 - Diagnose
 - Remedy

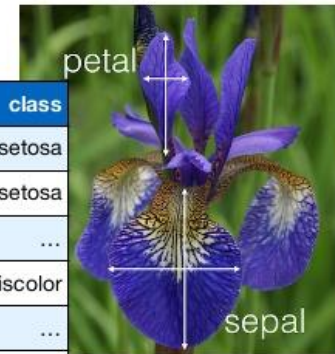
Nomenclature

IRIS

<https://archive.ics.uci.edu/ml/datasets/Iris>

Instances (samples, observations)

	sepal_length	sepal_width	petal_length	petal_width	class
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
...
50	6.4	3.2	4.5	1.5	vericolor
...
150	5.9	3.0	5.1	1.8	virginica



Features (attributes, dimensions)

Classes (targets)

$$\mathbf{X} = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1D} \\ x_{21} & x_{22} & \cdots & x_{2D} \\ x_{31} & x_{32} & \cdots & x_{3D} \\ \vdots & \vdots & \ddots & \vdots \\ \vdots & \vdots & \ddots & \vdots \\ x_{N1} & x_{N2} & \cdots & x_{ND} \end{bmatrix}$$

$$\mathbf{y} = [y_1, y_2, y_3, \cdots y_N]$$

<https://www.slideshare.net/SebastianRaschka/nextgen-talk-022015>

A simple example

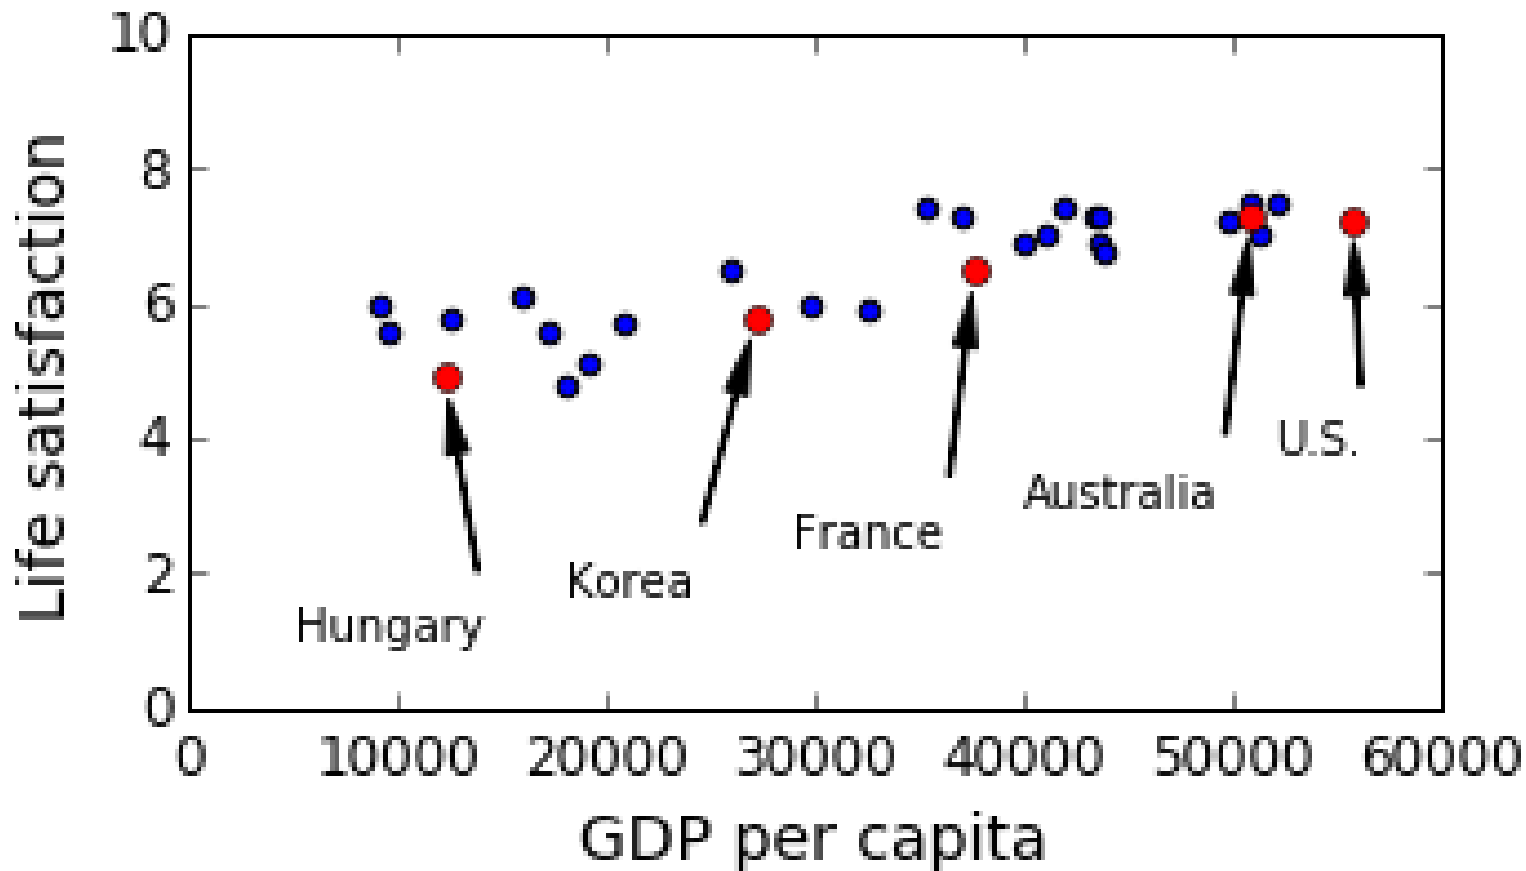


Figure from Aurelien Geron's ML book, page 19

Goal

- Learn a model from the data

Goal

- Build a model from the data

Goal

- **Train** a model from the data

Goal

- Train a model from the training data

Goal

- Train a model from the training data
- Make predictions

Training/Learning


- Which ones to choose?
- Or, a more general question is, How to train/learn these model parameters??
- The answer is by

defining a cost function
&
minimizing it

Cost function: basic idea

- Measures how bad (or good) a candidate model is
- Specifically, measure the difference between predictions from our model and the training data
- The objective is to minimize the cost function so as to minimize the difference between **predictions** and **observations**

Building a cost function



For simplicity, let us focus on one country, say, U.S.

- What is the **predicted value** for life satisfaction?
- What is **the value from training data**?

Difference between predicted and true values

- In our training data, we have **M** countries. Suppose U.S. is the i^{th} country.
- Predicted value: $h_{\theta}(x^{(i)})$
- True value: $y^{(i)}$
- Difference:

$$(h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- This is only for U.S.
- Remember that we want our model to fit all of our training data, not just one of them. Therefore, we sum the differences over all countries

$$\sum_{i=1}^M (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Cost function

$$\sum_{i=1}^M (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Remember

$$h_{\theta}(x^{(i)}) = \theta_0 + \theta_1 x^{(i)}$$

Cost function

$$\sum_{i=1}^M (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Remember

$$h_{\theta}(x^{(i)}) = \theta_0 + \theta_1 x^{(i)}$$

Therefore,

$$\sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$

Cost function

$$\sum_{i=1}^M (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Remember

$$h_{\theta}(x^{(i)}) = \theta_0 + \theta_1 x^{(i)}$$

Therefore,

$$J(\theta_0, \theta_1) = \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$

Minimization

- Cost function measures the difference between predicted and true values
- Remember that, we want to minimize this difference, i.e.,

$$\min J(\theta_0, \theta_1) = \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$

- Learning/training = Minimizing a cost function
- The process of learning a model from training data is essentially the process of minimizing a cost function.

Optimization

- Cost function measures the difference between predicted and true values
- Remember that, we want to minimize this difference, i.e.,

$$\min J(\theta_0, \theta_1) = \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$

- Learning/training = Minimizing a cost function
- The process of learning a model from training data is essentially the process of minimizing a cost function.
- Optimization: finding optimal parameter values that minimize a cost function

Optimization

- Cost function measures the difference between predicted and true values
- Remember that, we want to minimize this difference, i.e.,

$$\min J(\theta_0, \theta_1) = \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$

- Learning/training = Minimizing a cost function
- The process of learning a model from training data is essentially the process of **optimization**.
- **Optimization**: finding optimal parameter values that minimize a cost function

Best fit model

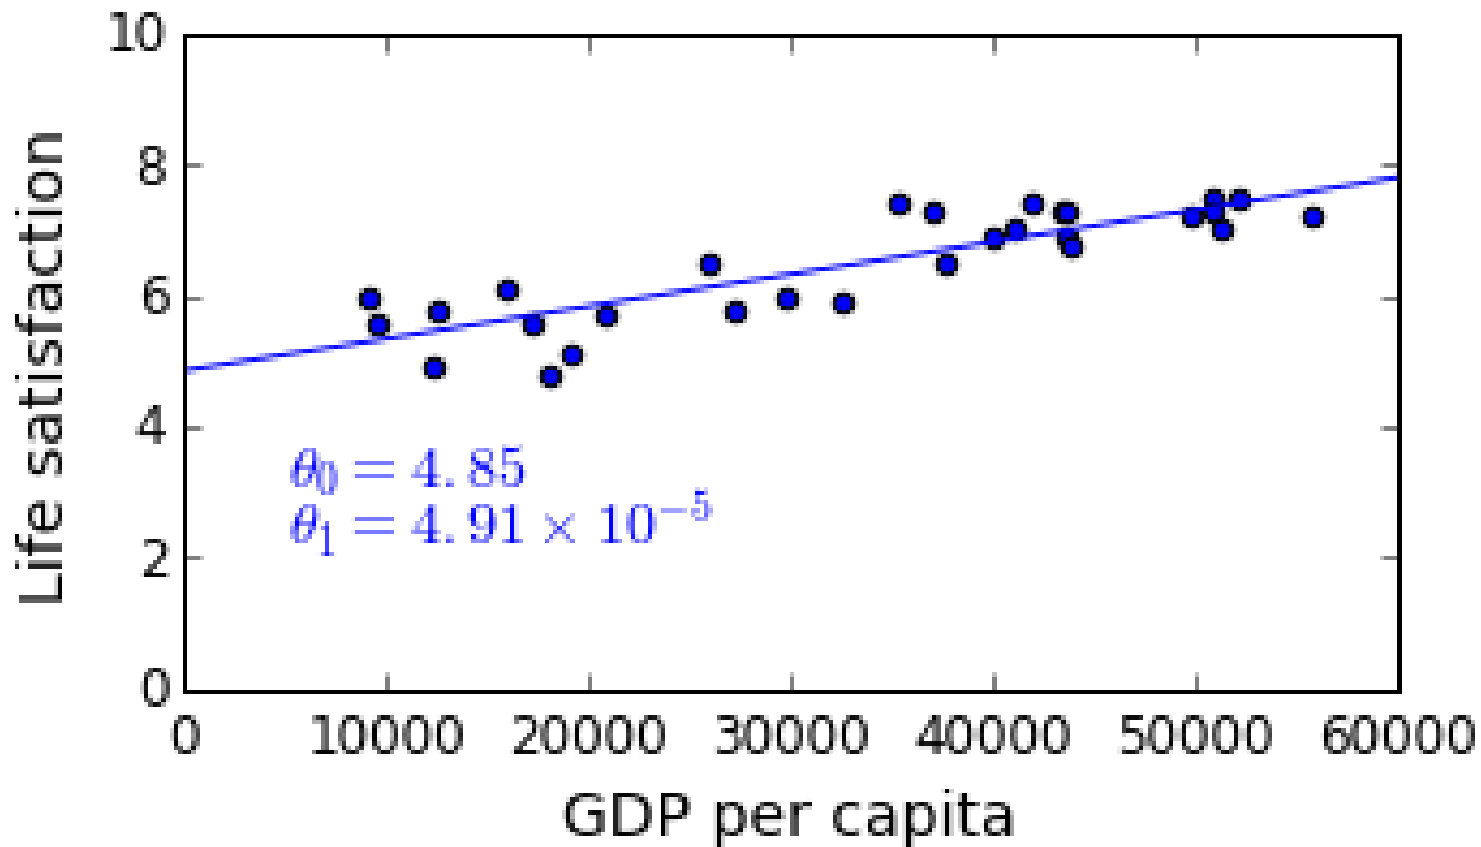


Figure from Aurelien Geron's ML book, page 20

Outline

- Machine learning basics
 - Example applications
 - Definitions
- Training
 - Cost function
 - Optimization
- Optimization algorithms
 - Batch gradient descent
 - Stochastic gradient descent
 - Mini-batch gradient descent
- Types of machine learning
 - Supervised vs. unsupervised
 - Regression vs. Classification
- Overfit vs Underfit
 - Diagnose
 - Remedy

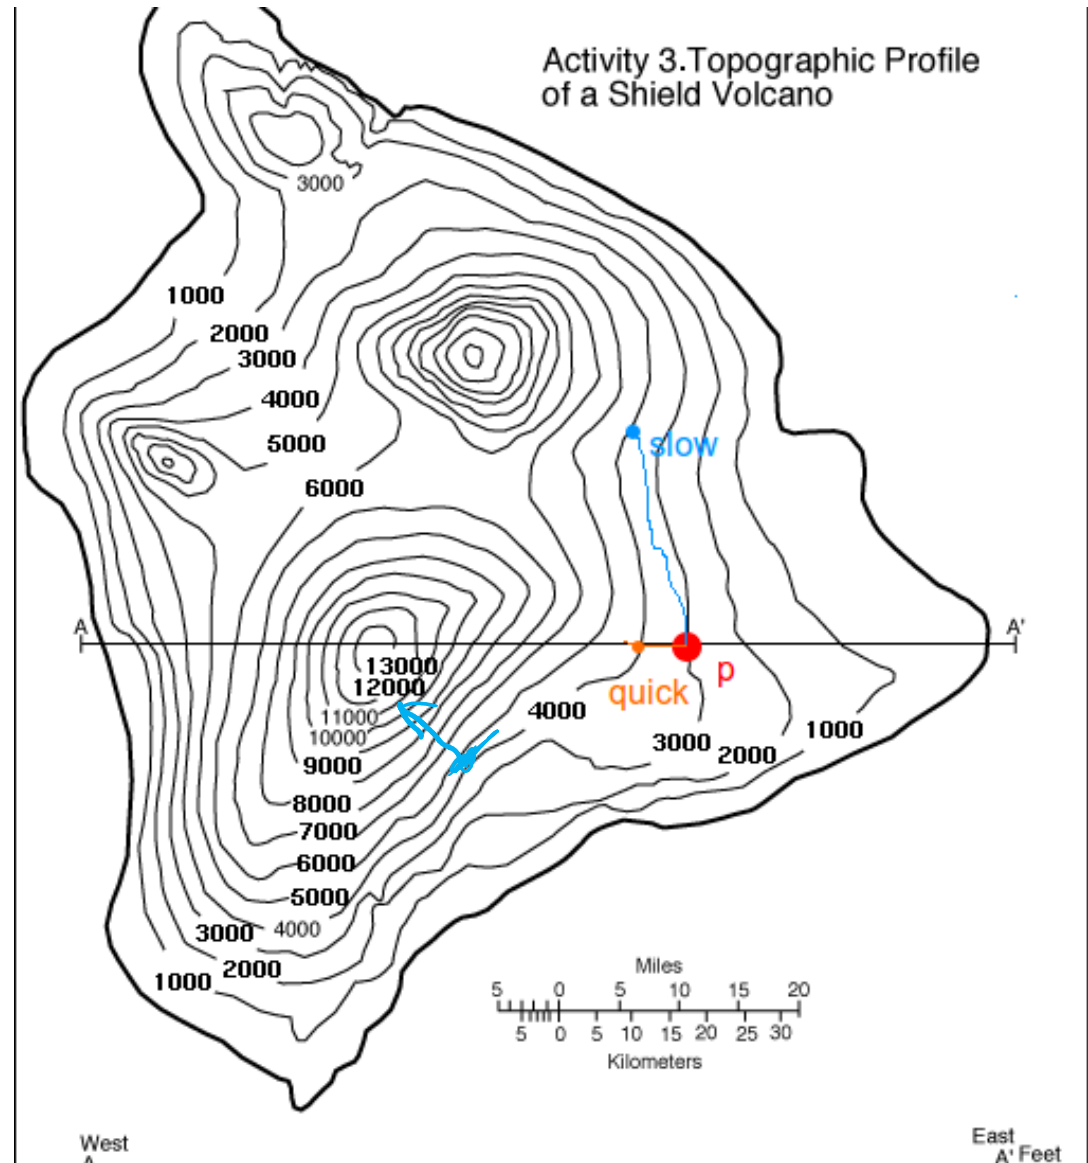
Gradient

- Gradient of a function $f(x, y)$ is defined as

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$$

Understanding gradient

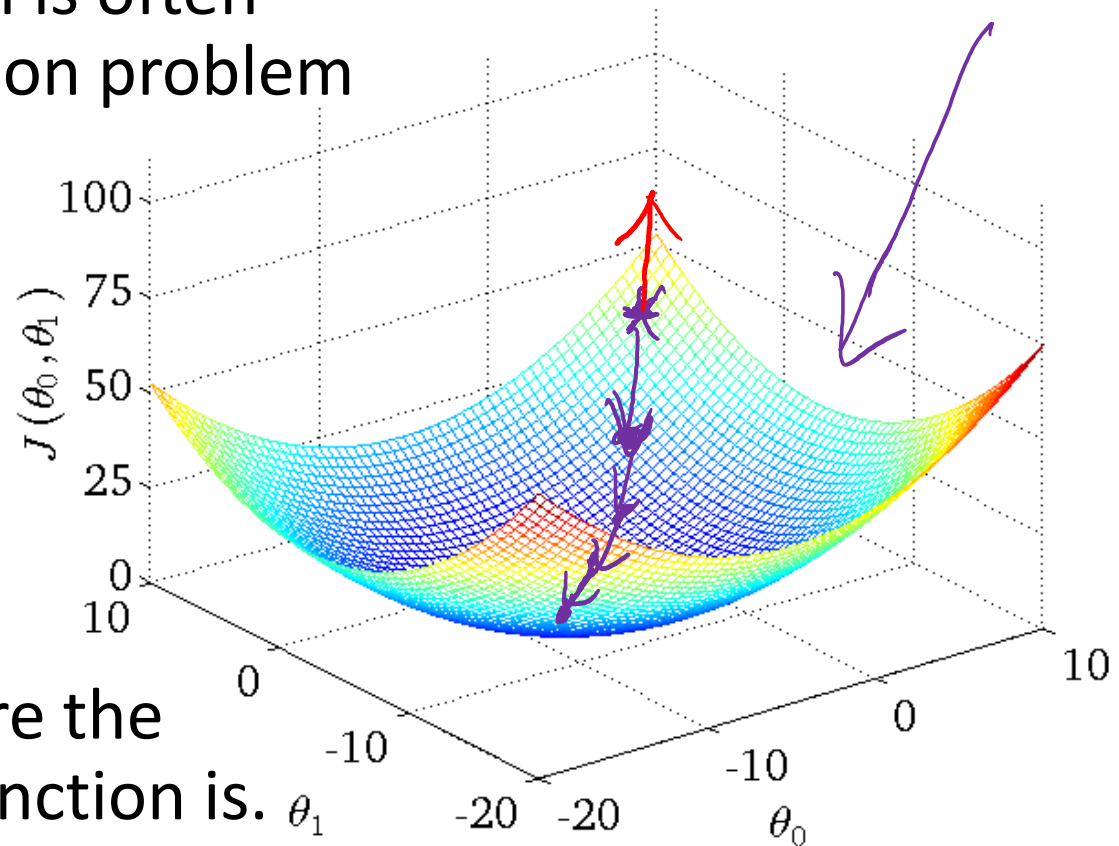
- Consider the topography as a 2D function $f(x, y)$
- The gradient direction tells you the fastest way up



Picture taken from <https://mathoverflow.net/questions/1977/why-is-the-gradient-normal>

Gradient in the context of optimization

- Optimization problem is often posed as a minimization problem



- We want to find where the minimum of a cost function is.

Picture taken from Andrew Ng's Machine Learning class on Coursera.org

Gradient descent algorithm

- Given initial values $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$
- While (not convergence):

$$\boldsymbol{\theta}^{(j)} = \boldsymbol{\theta}^{(j-1)} - \alpha \nabla J(\boldsymbol{\theta}^{(j-1)})$$

Gradient descent algorithm

- Given initial values $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$
- While (not convergence):

$$\theta_0^{(j)} = \theta_0^{(j-1)} - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0^{(j-1)}, \theta_1^{(j-1)})$$
$$\theta_1^{(j)} = \theta_1^{(j-1)} - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0^{(j-1)}, \theta_1^{(j-1)})$$

Gradient descent algorithm for linear regression

- Given initial values $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$
- While (not convergence):

$$\theta_0^{(j)} = \theta_0^{(j-1)} - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0^{(j-1)}, \theta_1^{(j-1)})$$
$$\theta_1^{(j)} = \theta_1^{(j-1)} - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0^{(j-1)}, \theta_1^{(j-1)})$$

$$\min J(\theta_0, \theta_1) = \frac{1}{2} \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$

Gradient descent algorithm for linear regression

- Given initial values $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$
- While (not convergence):

$$\theta_0^{(j)} = \theta_0^{(j-1)} - \alpha \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})$$

$$\theta_1^{(j)} = \theta_1^{(j-1)} - \alpha \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) x^{(i)}$$

- Each step of gradient descent uses ALL the training examples.

Batch gradient descent

- Each step of gradient descent uses ALL the training examples.

Problem with batch gradient descent

- When the number of training data is huge, say, $M = 300,000,000$, batch gradient descent becomes very slow.

Gradient descent algorithm

- Given initial values $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$
- While (not convergence):

$$\theta_0 = \theta_0 - \alpha \frac{1}{M} \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)})$$

$$\theta_1 = \theta_1 - \alpha \frac{1}{M} \sum_{i=1}^M (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) x^{(i)}$$

Stochastic gradient descent

- Given initial values $\boldsymbol{\theta}^{(0)} = [\theta_0^{(0)}, \theta_1^{(0)}]$

Repeat {

 Shuffle the training set

 For $i = 1, \dots, m$ {

$$\theta_0 = \theta_0 - \alpha (\theta_0 + \theta_1 x^{(i)} - y^{(i)})$$



$$\theta_1 = \theta_1 - \alpha (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) x^{(i)}$$

 }

}


SGD vs BGD

- SGD much faster
- Search path very irregular
- Cost function bounces up and down, decreasing only on average
- Over time, it ends up close to minimum, but never settles down.



Picture taken from
<https://www.cs.cmu.edu/~yuxiangw/docs/SSGD.pdf>

SGD cost function



<https://upload.wikimedia.org/wikipedia/commons/f/f3/Stogra.png>

Mini-batch gradient descent

- Mini-batch uses a small number ($1 < \# < M$) of training examples to update model parameters

Batch vs stochastic vs Mini-batch





Figure from Aurelien Geron's ML book, page 120


Outline

- Machine learning basics
 - Example applications
 - Definitions
- Training
 - Cost function
 - Optimization
- Optimization algorithms
 - Batch gradient descent
 - Stochastic gradient descent
 - Mini-batch gradient descent
- Types of machine learning
 - Supervised vs. unsupervised
 - Regression vs. Classification
- Overfit vs Underfit
 - Diagnose
 - Remedy

Supervised learning

- Training data come with labels (i.e., answers)
- Also termed labeled data
- E.g., cat classifier





Supervised learning: what is it?

- Let us consider each **image** as an input variable, x
- Also, consider each **label** as an output variable, y
- Supervised learning is all about learning a **mapping function** from the input to the output


$$y = f(x)$$

- So that, given a new image, x , your model learning model can predict y .


Unsupervised learning

- Training data does **not** have any label
- **Unlabeled data**
- The goal is to **discover the intrinsic, and often complicated structures** among data for better decision-making.
- **No answer available.** Algorithms are left to their own to discover the interesting structures in data.

Unsupervised learning: example




Unsupervised learning: example



Regression

- Predict **continuous numerical values**, such as **prices**, **temperatures**, etc.




Classification

- Predict **discrete categorical** values, such as class 1, 2, 3, etc.

Clas

- Pred 2, 3,



class 1,

Outline

- Machine learning basics
 - Example applications
 - Definitions
- Training
 - Cost function
 - Optimization
- Optimization algorithms
 - Batch gradient descent
 - Stochastic gradient descent
 - Mini-batch gradient descent
- Types of machine learning
 - Supervised vs. unsupervised
 - Regression vs. Classification
- Overfit vs Underfit
 - Diagnose
 - Remedy

Overfit: example




Figure from Aurelien Geron's ML book, page 19

Good fit




Figure from Aurelien Geron's ML book, page 20

Overfit (polynomial degree = 60)




Figure from Aurelien Geron's ML book, page 26

Overfit

- Fit the training data very well (actually too well)
- But, does not generalize well to new data.
- That is, predictions on new data will be bad!
- Remember, the whole purpose of machine learning is to make predictions.
- If a machine learning model only works well on training data, but not on new (i.e., unseen) data, it is NOT a good model/product.

How to tell if you are overfitting?

- Split your training data into three parts:
 1. Training set
 2. Validation set
 3. Test set

How to tell if you are overfitting?

- Split your training data into three parts:
 1. Training set
 2. Validation set
 3. Test set
- Use only training set for training (put the other two sets of data aside), calculate the prediction error J_{train}
- After training, apply the learned model to cross-validation set, calculate the prediction error J_{cv}
- If J_{train} is very small, J_{cv} is large, you overfit your data!

Remedy for overfitting

- **Overfitting** happens when your ML model is overly complex
- Therefore, **possible solutions** are:
 1. Collect more training data
 2. Reduce data noise
 3. Simplify model
 - using linear model rather than a high-degree polynomial model
 - using regularization
 - ...

Underfit

- The opposite of overfitting
- Your model is **too simple** to capture the meaningful information/structures/relations in the data.

Underfit: example





Figure from Aurelien Geron's ML book, page 121

Underfit: example



Overfit vs. underfit




Figure from Aurelien Geron's ML book, page 123

How to tell if you are underfitting?

- Split your training data into three parts:
 1. Training set
 2. Validation set
 3. Test set

How to tell if you are overfitting?

- Split your training data into three parts:
 1. Training set
 2. Validation set
 3. Test set
- Use only training set for training (put the other two sets of data aside), calculate the prediction error E^{train}
- After training, apply the learned model to cross-validation set, calculate the prediction error E^{cv}
- If E^{train} is large, E^{cv} is large, you underfit your data!

Remedy for underfitting

- Underfitting happens when your ML model is overly simple
- Therefore, possible solutions are:
 - ~~1. Collect more training data~~
 - ~~2. Reduce data noise~~
 3. Make your model more complex
 - using a high-degree polynomial model rather than a linear model
 - using less regularization
 - Adding more features such as (x_1^2, x_2^2, x_1x_2) to the learning algorithm (feature engineering)

Remember,

- If you are underfitting your data, collecting more data won't help!

Bias vs. Variances




Bias

- Due to over-simplified assumptions
- Your model is **heavily limited or biased** by your assumptions.
- E.g., assuming a linear model when the training data are actually from a non-linear model
- Lead to **underfitting** the training data

Variance


- Your model has **too much flexibility** (or, is allowed to have **too many small-scale variations**)
- E.g., assume a highly nonlinear model when the data are actually linear
- Lead to **overfitting** the data

Bias/variance



This slide is taken from Andrew Ng's ML class on coursera

Diagnosing bias vs. variance




Sample Question

- Suppose you trained a ML model on **training data**, and obtained a prediction accuracy of **89%** (i.e., prediction error of **11%**)
- Suppose you also applied the learned model to your **cross-validation data** set, and obtained a prediction accuracy of **96%** (i.e., a prediction error of **4%**)
- **Question**: Is this an **underfitting** or **overfitting** problem?

Remedy for overfitting

- **Overfitting** happens when your ML model is overly complex
- Therefore, **possible solutions** are:
 1. Collect more training data
 2. Reduce data noise
 3. Simplify model
 - using linear model rather than a high-degree polynomial model
 - using regularization
 - ...

Training data set of $M = 10$ points, each comprising an observation of input variable x and the corresponding target variable t .



The green curve shows the function $\sin(2\pi x)$ used to generate the data.

Goal: predict the value of t for some new value of x , based on the model learned from training data.

Polynomial curve fitting

Fit the data using a polynomial function of the form:

$$h_{\theta}(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \cdots + \theta_N x^N$$

where **N** is the degree of the polynomial

Polynomial curve fitting





Fit the data using a polynomial function of the form:

$$h_{\theta}(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \cdots + \theta_N x^N$$

where **N** is the degree of the polynomial


Question: how to select N?

Training models with $N = 0, 1, 3, 9$




Overfitting: the learned model captures noise, rather than the true and meaningful features/trends among the training data.

More training data




Learned models with $N = 9$ for $M = 10, 15$, and 100 .

More training data




Learned models with $N = 9$ for $M = 10, 15$, and 100.

More training data



Learned models with $N = 9$ for $M = 10, 15$, and 100 .


More training data



Learned models with $N = 9$ for $M = 10, 15$, and 100 .

Increasing the size of training data set helps decrease the overfitting.

More training data



Learned models with $N = 9$ for $M = 10, 15$, and 100 .

Increasing the size of training data set helps decrease the overfitting.

However, obtaining more training data is not always doable in reality.

Remedy for overfitting

- **Overfitting** happens when your ML model is overly complex
- Therefore, **possible solutions** are:
 1. Collect more training data
 2. Reduce data noise Not always feasible either!
 3. Simplify model
 - using linear model rather than a high-degree polynomial model
 - using regularization
 - ...

	N = 0	N = 1	N = 3	N = 9
θ_0	0.19	0.82	0.31	0.35
θ_1		-1.27	7.99	232.37
θ_2			-25.43	-5321.83
θ_3			17.37	48568.31
θ_4				-231639.30
θ_5				640042.26
θ_6				-1061800.52
θ_7				1042400.18
θ_8				-557682.99
θ_9				125201.43

Table of the coefficients θ learned from training data.

Note how the magnitudes of coefficients increase dramatically as the order of polynomial increases.

Regularization

- Discourage the learned model parameters (i.e., coefficients) from being too large.
- Keep the model simple.
- Keep the model from being unnecessarily complicated.

Regularization


$$J(\boldsymbol{\theta}) = \frac{1}{2} \sum_{i=1}^M (h_{\boldsymbol{\theta}}(x^{(i)}) - t^{(i)})^2$$

Regularization

$$J(\boldsymbol{\theta}) = \frac{1}{2} \sum_{i=1}^M (h_{\boldsymbol{\theta}}(x^{(i)}) - t^{(i)})^2 + \frac{1}{2} \lambda \sum_{j=1}^N \theta_j^2$$

Also known as **shrinkage** because it shrinks/reduces the values of the model parameters

Regularized curve fitting




$N = 9$ $M = 10$ with three different regularization parameters


How to select regularization parameter?

- Split the whole data set into **training** and **validation** data sets
- Carry out machine learning with different values for regularization parameters
- Calculate the **errors** for both **training** and **validation** data sets for each regularization parameter

Bias/variance as a function of regularization parameter



Diagnosing bias vs. variance



Regularization parameters

- Decision trees: `max_depth`

